

# Deep Learning

## Homework 1: Neural Networks

**Announced:** October 21 at 03:00 pm

**Deadline:** November 8 at 11:59 pm (eeclass)

**TA:**

蘇翁台 ([wengtai2008@hotmail.com](mailto:wengtai2008@hotmail.com))

林辰安 ([strike860930@gmail.com](mailto:strike860930@gmail.com))

陳冠宇 ([seed0123456@gmail.com](mailto:seed0123456@gmail.com))

---

### Introduction:

There is a saying that goes “Great oaks from little acorns grow.” In this assignment, you need to build a **shallow neural network** as well as the mini-batch SGD training process on MNIST dataset from scratch using Python3. The performance of your model should surpass 95% accuracy on testing data.

### Rules:

- **Built-in machine learning libraries (Sklearn, PyTorch, TensorFlow...) are not allowed to use for NN, you need to use basic mathematical operations in NumPy to define the behavior of each layer.**
- You can use any libraries to read the data from the MNIST folder and to do result visualization which should be shown on your report.
- Please properly comment your code to let us understand your train of thought.
- Discussions are encouraged, but plagiarism is strictly prohibited and punishable!

### Submission:

You should compress your code, dataset, report (hw1report\_studentID.pdf) and readme.txt (explain how to run your code) into a ZIP file (hw1\_studentID.zip), and submit it on **eeclass** before the due date. Zero credit for the submission after deadline.

### Implementation [80%, 10% for each part]

1. MNIST reading & splitting: (<http://yann.lecun.com/exdb/mnist/> for details) You need to split the original training data into validation data and training data. (proportion: 3/7)
2. Dense neural layer:  
Every output neuron has full connection to the input neurons.

3. ReLU layer:  
We add nonlinear activation functions after the neural layer using ReLU.
4. Softmax output:  
The final layer is typically a Softmax function which outputs the probability of a sample being in different classes.
5. Cross-entropy loss calculation:  
Use the output of a batch of data and their labels to calculate the CE loss.
6. Backward propagation:  
Propagate the error backwards and update each parameter.
7. Validation:  
Test the performance on validation data during each epoch.
8. Testing accuracy:  
**Reach 95% to get 10 points, 90~95% -> 7 points, 85~90% -> 5 points, zero otherwise.**

### **Report [20%]**

The format is not limited, but the following matters must be discussed in your report:

1. Show your model architecture and testing accuracy.
2. How do you implement feed forward and backward propagation? A brief explanation is fine.
3. Plot training loss and validation loss. (loss vs. epochs figure)
4. If we use a very deep NN with a large number of neurons, will the accuracy increase? Why or why not?
5. Why do we need to validate our model?
6. t-SNE results (**optional**, not included in 20% & **not need to implementation**):  
You can employ the function from: [[sklearn.manifold.TSNE](#)] or [[author page](#)]

### **Bonus [10%]**

Try to use t-SNE (read the document before you call function) to visualize the outputs of (1) validation data at different epochs and (2) testing data, then show the result on your report.